



DralFire / ProcessBigData25B

[Code](#)[Issues](#)[Pull requests](#)[Actions](#)[Projects](#)[Wiki](#)[Security](#)[In](#)

ProcessBigData25B / 00 - Challenge 3

/ Challenge_3_HeartDisease_SparkML_v2 - PEREZ ROSAS LUIS ALFREDO.ipynb [🔗](#)



DralFire Challenge 3 Avanzado

c65cb05 · 3 minutes ago



855 lines (855 loc) · 194 KB

Challenge 3 Avanzado

Enfermedades Cardiovasculares ❤️

PÉREZ ROSAS LUIS ALFREDO | MCD | GRUPO 2 | UDG

Este challenge avanzado permitirá a los estudiantes aplicar PySpark para analizar grandes volúmenes de datos relacionados con un tema cercano a su investigación de tesis, integrando los conocimientos adquiridos en los challenges anteriores.

Utilizar APIs avanzadas como Spark Streaming o Spark ML para trabajar con una gran base de datos en tiempo real o construir un modelo predictivo complejo.

1) Parámetros de ejecución

Ajusta la ruta al CSV (usa `,` como separador por defecto). Si tu archivo tiene otra columna objetivo distinta a `target` o `num`, ajusta `LABEL_BASE_COL`.

In [2]:

```
# Parámetros
DATA_PATH = "data/heart_disease_uci.csv" # <-- CAMBIA ESTA RUTA
CSV_SEP = "," # separador común del CSV
OUTDIR = "outputs_heart" # carpeta de salida

# Si conoces La columna objetivo exacta, indícalo aquí (por defecto intentará
LABEL_BASE_COL = None # e.g., "target" o "num"; si No

# %matplotlib inline # Descomenta si estás en Jupyter clásico
```

2) Inicializar Spark

In [3]:

```
from pyspark.sql import SparkSession

spark = (
    SparkSession.builder
    .appName("challenge-3-avanzado-heart-v2")
    .getOrCreate()
)
spark
```

Out[3]: **SparkSession - in-memory**

SparkContext

[Spark UI](#)

Version v3.5.7
Master local[*]
AppName challenge-3-avanzado-heart-v2

3) Carga de datos

```
In [4]: from pyspark.sql import functions as F
import os

os.makedirs(OUTDIR, exist_ok=True)

raw = spark.read.csv(DATA_PATH, header=True, inferSchema=True, sep=CSV_SEP)
raw.printSchema()
raw.show(5)
```

```
root
 |-- id: integer (nullable = true)
 |-- age: integer (nullable = true)
 |-- sex: string (nullable = true)
 |-- dataset: string (nullable = true)
 |-- cp: string (nullable = true)
 |-- trestbps: integer (nullable = true)
 |-- chol: integer (nullable = true)
 |-- fbs: boolean (nullable = true)
 |-- restecg: string (nullable = true)
 |-- thalch: integer (nullable = true)
 |-- exang: boolean (nullable = true)
 |-- oldpeak: double (nullable = true)
 |-- slope: string (nullable = true)
 |-- ca: integer (nullable = true)
 |-- thal: string (nullable = true)
 |-- num: integer (nullable = true)

+-----+-----+-----+-----+
| id|age| sex| dataset|          cp|trestbps|chol|   fbs|      restecg|tha
lch|exang|oldpeak|       slope| ca|          thal|num|
+-----+-----+-----+-----+
|  1| 63| Male|Cleveland| typical angina|    145| 233| true|lv hypertrophy|
150|false|    2.3|downsloping|  0| fixed defect|  0|
|  2| 67| Male|Cleveland| asymptomatic|   160| 286|false|lv hypertrophy|
108| true|    1.5|     flat|  3| normal|  2|
|  3| 67| Male|Cleveland| asymptomatic|   120| 229|false|lv hypertrophy|
129| true|    2.6|     flat|  2|reversible defect|  1|
|  4| 37| Male|Cleveland| non-anginal|   130| 250|false|      normal|
187|false|    3.5|downsloping|  0| normal|  0|
|  5| 41|Female|Cleveland|atypical angina|   130| 204|false|lv hypertrophy|
172|false|    1.4| upsloping|  0| normal|  0|
+-----+-----+-----+-----+
-----+-----+-----+-----+
only showing top 5 rows
```

Diccionario de Datos

id :: Identificador del paciente.
 age :: Edad.
 gender :: Género.
 height :: Estatura.
 weight :: Peso.
 ap_hi :: Presión arterial sistólica
 ap_lo :: Presión arterial diastólica
 cholesterol :: Nivel de colesterol
 gluc :: Nivel de glucosa:
 smoke :: ¿Es fumador?
 alco :: ¿Bebe alcohol?
 active :: ¿Realiza actividad física?
 cardio :: ¿Tiene alguna enfermedad cardiovascular diagnosticada?

4) Detección de la columna objetivo y creación de label binaria

In [5]:

```

# Detectar columna objetivo si no se especificó
candidates = [c for c in raw.columns if c.lower() in ("target", "num")]
base_col = LABEL_BASE_COL if LABEL_BASE_COL else (candidates[0] if candidates

if base_col is None:
    raise ValueError("No se pudo detectar la columna objetivo automáticamente.
                      "Indica LABEL_BASE_COL con el nombre correcto (p. ej., 't

print("Columna base para etiqueta:", base_col)

# Mapeo a binaria: Label = 1 si valor > 0, 0 en otro caso
df = raw.withColumn("label", (F.col(base_col).cast("double") > F.lit(0.0)).cast("integer"))
df.select(base_col, "label").groupBy("label").count().orderBy("label").show()
  
```

```

Columna base para etiqueta: num
+----+-----+
|label|count|
+----+-----+
|    0|  411|
|    1|  509|
+----+-----+
  
```

5) EDA con gráficas

Importante: Usamos `matplotlib` (una figura por gráfico, sin estilos ni colores específicos).

Se convertirán los datos a pandas con `toPandas()` para graficar. Si el dataset es muy grande, considera muestrear (`sample()`).

In [6]:

```

import matplotlib.pyplot as plt

# Seleccionar solo columnas numéricas para las gráficas
numeric_cols = [c for (c, t) in df.dtypes if t in ("int", "bigint", "double",
pdf = df.select(numeric_cols).toPandas()


# 5.1 Histogramas
ax = pdf.hist(figsize=(10, 8))
plt.tight_layout()
plt.savefig(os.path.join(OUTDIR, "histogramas.png"))
plt.show()


# 5.2 Boxplots
plt.figure()
pdf.boxplot(rot=90)
plt.title("Boxplots de variables numéricas")
plt.tight_layout()
plt.savefig(os.path.join(OUTDIR, "boxplots.png"))
plt.show()

# 5.3 Matriz de correlación
plt.figure()
corr = pdf.corr(numeric_only=True)
plt.imshow(corr, interpolation='nearest')
plt.xticks(range(len(corr.columns)), corr.columns, rotation=90)
plt.yticks(range(len(corr.columns)), corr.columns)
plt.colorbar()
plt.title("Matriz de correlación")
plt.tight_layout()
plt.savefig(os.path.join(OUTDIR, "correlacion.png"))
plt.show()


# 5.4 Distribución de La etiqueta
if "label" in df.columns:
    label_pdf = df.select("label").toPandas()
    plt.figure()
    label_pdf["label"].value_counts().sort_index().plot(kind="bar")
    plt.title("Distribución de la etiqueta")
    plt.xlabel("label")
    plt.ylabel("conteo")
    plt.tight_layout()
    plt.savefig(os.path.join(OUTDIR, "distribucion_label.png"))
    plt.show()

```







Boxplots de variables numéricas




Matriz de correlación





Distribución de la etiqueta



6) Preparación de datos y *features*

In [7]:

```
# Revisa tu esquema
df.printSchema()

# Columnas del ensamblador (ajusta si usaste otro nombre)
# Evita meter 'label' ni la columna objetivo original (target/num)
exclude = {"label", "target", "num", "features", "features_raw"}
feature_cols = [c for c, t in df.dtypes if c not in exclude]

# Qué columnas no son numéricas aún
non_numeric = [c for c, t in df.dtypes if c in feature_cols and t not in ("int", "double")]
missing = [c for c in feature_cols if c not in dict(df.dtypes)]
print("No numéricas:", non_numeric)
print("Inexistentes:", missing)
```

```
root
|-- id: integer (nullable = true)
|-- age: integer (nullable = true)
|-- sex: string (nullable = true)
```

ProcessBigData25B/00 - Challenge 3/Challenge_3_HeartDisease_SparkML_v2 - PEREZ ROSAS LUIS ALFREDO.ipynb at main · Dra...

```

    |-- dataset: string (nullable = true)
    |-- cp: string (nullable = true)
    |-- trestbps: integer (nullable = true)
    |-- chol: integer (nullable = true)
    |-- fbs: boolean (nullable = true)
    |-- restecg: string (nullable = true)
    |-- thalch: integer (nullable = true)
    |-- exang: boolean (nullable = true)
    |-- oldpeak: double (nullable = true)
    |-- slope: string (nullable = true)
    |-- ca: integer (nullable = true)
    |-- thal: string (nullable = true)
    |-- num: integer (nullable = true)
    |-- label: integer (nullable = true)

```

No numéricas: ['sex', 'dataset', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'thal']
 Inexistentes: []

In [24]:

```

from pyspark.sql.functions import col, regexp_replace, trim
from pyspark.sql import functions as F

df_cast = df

# Elimina filas con nulos en cualquier columna
df_cast.na.drop()

for c in non_numeric:
    # Limpia espacios y comas decimales si las hubiera
    df_cast = df_cast.withColumn(c, trim(col(c)))
    df_cast = df_cast.withColumn(c, regexp_replace(col(c), ",", "."))
    df_cast = df_cast.withColumn(c, col(c).cast("double"))

# Recalcula solo columnas numéricas reales para el assembler
numeric_feature_cols = [
    c for c, t in df_cast.dtypes
    if c in feature_cols and t in ("int", "bigint", "double", "float")
]
print("Usaré estas columnas numéricas:", numeric_feature_cols)

print(df_cast)

```

Usaré estas columnas numéricas: ['id', 'age', 'sex', 'dataset', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalch', 'exang', 'oldpeak', 'slope', 'ca', 'thal']

DataFrame[id: int, age: int, sex: double, dataset: double, cp: double, trestbps: int, chol: int, fbs: double, restecg: double, thalch: int, exang: double, oldpeak: double, slope: double, ca: int, thal: double, num: int, label: int]

In [25]:

```

from pyspark.ml.feature import VectorAssembler, StandardScaler
from pyspark.ml import Pipeline

```

In [26]:

```

#
#assembler = VectorAssembler(
#    inputCols=numeric_feature_cols,
#    outputCol="features raw"

```

```
ProcessBigData25B/00 - Challenge 3/Challenge_3_HeartDisease_SparkML_v2 - PEREZ ROSAS LUIS ALFREDO.ipynb at main · Dra...
```

```

#     handleInvalid="skip" # en Spark 3.5 soportado; salta filas con nulos/in
#)
#scaler = StandardScaler(inputCol="features_raw", outputCol="features")

#prep = Pipeline(stages=[assembler, scaler])

assembler = VectorAssembler(inputCols=numeric_feature_cols, outputCol="feature"
scaler = StandardScaler(inputCol="features_raw", outputCol="features")
prep = Pipeline(stages=[assembler, scaler])

prepared = prep.fit(df_cast).transform(df_cast)

train, test = prepared.randomSplit([0.8, 0.2], seed=42)
train.count(), test.count()

```

Out[26]: (771, 149)

In []:

7) Modelado: Logistic Regression y Random Forest

In [40]:

```

from pyspark.ml.classification import LogisticRegression, RandomForestClassifier
from pyspark.ml.feature import Imputer

# Impute missing values in numeric features before modeling
imputer = Imputer(
    inputCols=numeric_feature_cols,
    outputCols=[f"{c}_imputed" for c in numeric_feature_cols]
).setStrategy("mean")

df_cast.na.drop()
df_cast.dropna()
df_cast.dropna(how='all')

df_imputed = imputer.fit(df_cast).transform(df_cast)

# Update assembler to use imputed columns
imputed_feature_cols = [f"{c}_imputed" for c in numeric_feature_cols]
assembler = VectorAssembler(inputCols=imputed_feature_cols, outputCol="feature"
scaler = StandardScaler(inputCol="features_raw", outputCol="features")
prep = Pipeline(stages=[assembler, scaler])

prepared = prep.fit(df_imputed).transform(df_imputed)
train, test = prepared.randomSplit([0.8, 0.2], seed=42)

lr = LogisticRegression(featuresCol="features", labelCol="label")
rf = RandomForestClassifier(featuresCol="features", labelCol="label", numTrees

lr_model = lr.fit(train)
rf_model = rf.fit(train)

pred_lr = lr_model.transform(test)

```

```

pred_rf = rf_model.transform(test)

pred_lr.select("label", "prediction", "probability").show(5, truncate=False)
pred_rf.select("label", "prediction", "probability").show(5, truncate=False)

```

Py4JJavaError

Cell In[40], line 14

```

11 df_cast.dropna()
12 df_cast.dropna(how='all')
--> 14 df_imputed = imputer.fit(df_cast).transform(df_cast)
15 # Update assembler to use imputed columns
17 imputed_feature_cols = [f"${c}_imputed" for c in numeric_feature_cols]

```

File c:\Users\alfre\anaconda3\envs\ProcessBigData25B\lib\site-packages\pyspark\m
1\base.py:205, in Estimator.fit(self, dataset, params)

```

203         return self.copy(params)._fit(dataset)
204     else:
--> 205         return self._fit(dataset)
206 else:
207     raise TypeError(
208         "Params must be either a param map or a list/tuple of param map
s, "
209         "but got %s." % type(params)
210     )

```

File c:\Users\alfre\anaconda3\envs\ProcessBigData25B\lib\site-packages\pyspark\m
1\wrapper.py:381, in JavaEstimator._fit(self, dataset)

```

380 def _fit(self, dataset: DataFrame) -> JM:
--> 381     java_model = self._fit_java(dataset)
382     model = self._create_model(java_model)
383     return self._copyValues(model)

```

File c:\Users\alfre\anaconda3\envs\ProcessBigData25B\lib\site-packages\pyspark\m
1\wrapper.py:378, in JavaEstimator._fit_java(self, dataset)

```

375 assert self._java_obj is not None
377 self._transfer_params_to_java()
--> 378 return self._java_obj.fit(dataset._jdf)

```

File c:\Users\alfre\anaconda3\envs\ProcessBigData25B\lib\site-packages\py4j\java
_gateway.py:1322, in JavaMember.__call__(self, *args)

```

1316 command = proto.CALL_COMMAND_NAME +\
1317     self.command_header +\
1318     args_command +\
1319     proto.END_COMMAND_PART
1321 answer = self.gateway_client.send_command(command)
-> 1322 return_value = get_return_value(
1323     answer, self.gateway_client, self.target_id, self.name)
1325 for temp_arg in temp_args:
1326     if hasattr(temp_arg, "_detach"):

```

File c:\Users\alfre\anaconda3\envs\ProcessBigData25B\lib\site-packages\pyspark\er
rors\exceptions\captured.py:179, in capture_sql_exception.<locals>.deco(*a, **k
w)

```

177 def deco(*a: Any, **kw: Any) -> Any:
178     try:
--> 179         return f(*a, **kw)
180     except Py4JJavaError as e:
181         converted = convert_exception(e.java_exception)

```

```

File c:\Users\alfre\anaconda3\envs\ProcessBigData25B\lib\site-packages\py4j\proto
ocol.py:326, in get_return_value(answer, gateway_client, target_id, name)
    324 value = OUTPUT_CONVERTER[type](answer[2:], gateway_client)
    325 if answer[1] == REFERENCE_TYPE:
--> 326     raise Py4JJavaError(
    327         "An error occurred while calling {0}{1}{2}.\\n".
    328         format(target_id, ".", name), value)
    329 else:
    330     raise Py4JError(
    331         "An error occurred while calling {0}{1}{2}. Trace:\\n{3}\\n".
    332         format(target_id, ".", name, value))

```

Py4JJavaError: An error occurred while calling o1244.fit.
: org.apache.spark.SparkException: surrogate cannot be computed. All the values
in sex,dataset,cp,fbs,restecg,exang,slope,thal are Null, Nan or missingValue(Na
N)
at org.apache.spark.ml.feature.Imputer.fit(Imputer.scala:199)
at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke0(Nativ
e Method)
at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke(Native
MethodAccessorImpl.java:77)
at java.base/jdk.internal.reflect.DelegatingMethodAccessorImpl.invoke(De
legatingMethodAccessorImpl.java:43)
at java.base/java.lang.reflect.Method.invoke(Method.java:568)
at py4j.reflection.MethodInvoker.invoke(MethodInvoker.java:244)
at py4j.reflection.ReflectionEngine.invoke(ReflectionEngine.java:374)
at py4j.Gateway.invoke(Gateway.java:282)
at py4j.commands.AbstractCommand.invokeMethod(AbstractCommand.java:132)
at py4j.commands.CallCommand.execute(CallCommand.java:79)
at py4j.ClientServerConnection.waitForCommands(ClientServerConnection.ja
va:182)
at py4j.ClientServerConnection.run(ClientServerConnection.java:106)
at java.base/java.lang.Thread.run(Thread.java:842)

8) Evaluación (AUC)

In [32]:

```

from pyspark.ml.evaluation import BinaryClassificationEvaluator

evaluator = BinaryClassificationEvaluator(labelCol="label", rawPredictionCol="rawPrediction")
auc_lr = evaluator.evaluate(pred_lr)
auc_rf = evaluator.evaluate(pred_rf)

print("AUC LR:", auc_lr)
print("AUC RF:", auc_rf)

```

NameError Traceback (most recent call last)
Cell In[32], line 4
1 from pyspark.ml.evaluation import BinaryClassificationEvaluator
3 evaluator = BinaryClassificationEvaluator(labelCol="label", rawPredictio
nCol="rawPrediction")
----> 4 auc_lr = evaluator.evaluate(pred_lr)
5 auc_rf = evaluator.evaluate(pred_rf)
7 print("AUC LR:", auc_lr)

NameError: name 'pred_lr' is not defined

9) Curvas ROC y PR (LR, training summary)

In [33]:

```
# Curvas obtenidas del summary de entrenamiento de LR
try:
    train_summary = lr_model.summary

    # ROC
    roc_df = train_summary.roc.toPandas()
    plt.figure()
    plt.plot(roc_df['FPR'], roc_df['TPR'])
    plt.xlabel('FPR')
    plt.ylabel('TPR')
    plt.title('Curva ROC (LR, training)')
    plt.tight_layout()
    plt.savefig(os.path.join(OUTDIR, "roc_lr_training.png"))
    plt.show()

    # PR
    pr_df = train_summary.pr.toPandas()
    plt.figure()
    plt.plot(pr_df['recall'], pr_df['precision'])
    plt.xlabel('Recall')
    plt.ylabel('Precision')
    plt.title('Curva Precision-Recall (LR, training)')
    plt.tight_layout()
    plt.savefig(os.path.join(OUTDIR, "pr_lr_training.png"))
    plt.show()
except Exception as e:
    print("[WARN] No se pudieron generar curvas LR:", e)
```

[WARN] No se pudieron generar curvas LR: name 'lr_model' is not defined

10) Matrices de confusión (Spark puro)

In [34]:

```
import numpy as np

def confusion_matrix_spark(pred_df):
    cm_pdf = (
        pred_df.groupBy("label", "prediction")
        .count()
        .toPandas()
    )
    labels = sorted(set(cm_pdf["label"]).union(set(cm_pdf["prediction"])))
    size = len(labels)
    mat = [[0 for _ in range(size)] for _ in range(size)]
    for _, row in cm_pdf.iterrows():
        i = labels.index(row["label"])
        j = labels.index(row["prediction"])
        mat[i][j] = int(row["count"])
    return np.array(mat), labels
```

```

plt.ylabel("Real")
for i in range(cm_rf.shape[0]):
    for j in range(cm_rf.shape[1]):
        plt.text(j, i, int(cm_rf[i, j]), ha="center", va="center")
plt.colorbar()
plt.tight_layout()
plt.savefig(os.path.join(OUTDIR, "confusion_rf.png"))
plt.show()

cm_lr, labels_lr = confusion_matrix_spark(pred_lr)
plt.figure()
plt.imshow(cm_lr, interpolation="nearest")
plt.title("Matriz de confusión (LR)")
plt.xlabel("Predicho")
plt.ylabel("Real")
for i in range(cm_lr.shape[0]):
    for j in range(cm_lr.shape[1]):
        plt.text(j, i, int(cm_lr[i, j]), ha="center", va="center")
plt.colorbar()
plt.tight_layout()
plt.savefig(os.path.join(OUTDIR, "confusion_lr.png"))
plt.show()

```

NameError Traceback (most recent call last)
Cell In[34], line 18
15 mat[i][j] = int(row["count"])
16 return np.array(mat), labels
---> 18 cm_rf, labels_rf = confusion_matrix_spark(pred_rf)
19 plt.figure()
20 plt.imshow(cm_rf, interpolation="nearest")

NameError: name 'pred_rf' is not defined

11) Importancia de variables (RF)

In [35]:

```

import pandas as pd

fi = pd.DataFrame({
    "feature": feature_cols,
    "importance": rf_model.featureImportances.toArray()
}).sort_values("importance", ascending=False)

fi.head(15)

```

NameError Traceback (most recent call last)
Cell In[35], line 5
1 import pandas as pd
3 fi = pd.DataFrame({
4 "feature": feature_cols,
---> 5 "importance": rf_model.featureImportances.toArray()
6 }).sort_values("importance", ascending=False)
8 fi.head(15)

NameError: name 'rf_model' is not defined

In [15]:

```
# Gráfico de importancia
plt.figure()
plt.bar(fi["feature"], fi["importance"])
plt.xticks(rotation=90)
plt.title("Importancia de variables (RF) – Heart Disease")
plt.tight_layout()
plt.savefig(os.path.join(OUTDIR, "feature_importance_rf.png"))
plt.show()

fi.to_csv(os.path.join(OUTDIR, "feature_importance_rf.csv"), index=False)
```

NameError

Traceback (most recent call last)

Cell In[15], line 3

```
1 # Gráfico de importancia
2 plt.figure()
----> 3 plt.bar(fi["feature"], fi["importance"])
4 plt.xticks(rotation=90)
5 plt.title("Importancia de variables (RF) – Heart Disease")
```

NameError: name 'fi' is not defined
<Figure size 640x480 with 0 Axes>

12) Conclusiones